

reinforcement_learning_advanced_ant_walk

May 2, 2020

1 Twin-Delayed DDPG

1.1 Installing the packages

```
[0]: !pip install pybullet
```

1.2 Importing the libraries

```
[0]: import os
import time
import random
import numpy as np
import matplotlib.pyplot as plt
import pybullet_envs
import gym
import torch
import torch.nn as nn
import torch.nn.functional as F
from gym import wrappers
from torch.autograd import Variable
from collections import deque
```

1.3 Step 1: We initialize the Experience Replay memory

```
[0]: class ReplayBuffer(object):

    def __init__(self, max_size=1e6):
        self.storage = []
        self.max_size = max_size
        self.ptr = 0

    def add(self, transition):
        if len(self.storage) == self.max_size:
            self.storage[int(self.ptr)] = transition
```

```

        self.ptr = (self.ptr + 1) % self.max_size
    else:
        self.storage.append(transition)

def sample(self, batch_size):
    ind = np.random.randint(0, len(self.storage), size=batch_size)
    batch_states, batch_next_states, batch_actions, batch_rewards, batch_dones =
    → [], [], [], [], []
    for i in ind:
        state, next_state, action, reward, done = self.storage[i]
        batch_states.append(np.array(state, copy=False))
        batch_next_states.append(np.array(next_state, copy=False))
        batch_actions.append(np.array(action, copy=False))
        batch_rewards.append(np.array(reward, copy=False))
        batch_dones.append(np.array(done, copy=False))
    return np.array(batch_states), np.array(batch_next_states), np.
    → array(batch_actions), np.array(batch_rewards).reshape(-1, 1), np.
    → array(batch_dones).reshape(-1, 1)

```

1.4 Step 2: We build one neural network for the Actor model and one neural network for the Actor target

```

[0]: class Actor(nn.Module):

    def __init__(self, state_dim, action_dim, max_action):
        super(Actor, self).__init__()
        self.layer_1 = nn.Linear(state_dim, 400)
        self.layer_2 = nn.Linear(400, 300)
        self.layer_3 = nn.Linear(300, action_dim)
        self.max_action = max_action

    def forward(self, x):
        x = F.relu(self.layer_1(x))
        x = F.relu(self.layer_2(x))
        x = self.max_action * torch.tanh(self.layer_3(x))
        return x

```

1.5 Step 3: We build two neural networks for the two Critic models and two neural networks for the two Critic targets

```

[0]: class Critic(nn.Module):

    def __init__(self, state_dim, action_dim):
        super(Critic, self).__init__()

```

```

# Defining the first Critic neural network
self.layer_1 = nn.Linear(state_dim + action_dim, 400)
self.layer_2 = nn.Linear(400, 300)
self.layer_3 = nn.Linear(300, 1)
# Defining the second Critic neural network
self.layer_4 = nn.Linear(state_dim + action_dim, 400)
self.layer_5 = nn.Linear(400, 300)
self.layer_6 = nn.Linear(300, 1)

def forward(self, x, u):
    xu = torch.cat([x, u], 1)
    # Forward-Propagation on the first Critic Neural Network
    x1 = F.relu(self.layer_1(xu))
    x1 = F.relu(self.layer_2(x1))
    x1 = self.layer_3(x1)
    # Forward-Propagation on the second Critic Neural Network
    x2 = F.relu(self.layer_4(xu))
    x2 = F.relu(self.layer_5(x2))
    x2 = self.layer_6(x2)
    return x1, x2

def Q1(self, x, u):
    xu = torch.cat([x, u], 1)
    x1 = F.relu(self.layer_1(xu))
    x1 = F.relu(self.layer_2(x1))
    x1 = self.layer_3(x1)
    return x1

```

1.6 Steps 4 to 15: Training Process

```

[0]: # Selecting the device (CPU or GPU)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Building the whole Training Process into a class

class TD3(object):

    def __init__(self, state_dim, action_dim, max_action):
        self.actor = Actor(state_dim, action_dim, max_action).to(device)
        self.actor_target = Actor(state_dim, action_dim, max_action).to(device)
        self.actor_target.load_state_dict(self.actor.state_dict())
        self.actor_optimizer = torch.optim.Adam(self.actor.parameters())
        self.critic = Critic(state_dim, action_dim).to(device)
        self.critic_target = Critic(state_dim, action_dim).to(device)
        self.critic_target.load_state_dict(self.critic.state_dict())
        self.critic_optimizer = torch.optim.Adam(self.critic.parameters())

```

```

self.max_action = max_action

def select_action(self, state):
    state = torch.Tensor(state.reshape(1, -1)).to(device)
    return self.actor(state).cpu().data.numpy().flatten()

def train(self, replay_buffer, iterations, batch_size=100, discount=0.99,
→tau=0.005, policy_noise=0.2, noise_clip=0.5, policy_freq=2):

    for it in range(iterations):

        # Step 4: We sample a batch of transitions (s, s', a, r) from the memory
        batch_states, batch_next_states, batch_actions, batch_rewards,
→batch_dones = replay_buffer.sample(batch_size)
        state = torch.Tensor(batch_states).to(device)
        next_state = torch.Tensor(batch_next_states).to(device)
        action = torch.Tensor(batch_actions).to(device)
        reward = torch.Tensor(batch_rewards).to(device)
        done = torch.Tensor(batch_dones).to(device)

        # Step 5: From the next state s', the Actor target plays the next action
→a'
        next_action = self.actor_target(next_state)

        # Step 6: We add Gaussian noise to this next action a' and we clamp it in
→a range of values supported by the environment
        noise = torch.Tensor(batch_actions).data.normal_(0, policy_noise).
→to(device)
        noise = noise.clamp(-noise_clip, noise_clip)
        next_action = (next_action + noise).clamp(-self.max_action, self.
→max_action)

        # Step 7: The two Critic targets take each the couple (s', a') as input
→and return two Q-values Qt1(s',a') and Qt2(s',a') as outputs
        target_Q1, target_Q2 = self.critic_target(next_state, next_action)

        # Step 8: We keep the minimum of these two Q-values: min(Qt1, Qt2)
        target_Q = torch.min(target_Q1, target_Q2)

        # Step 9: We get the final target of the two Critic models, which is: Qt
→= r + * min(Qt1, Qt2), where is the discount factor
        target_Q = reward + ((1 - done) * discount * target_Q).detach()

        # Step 10: The two Critic models take each the couple (s, a) as input and
→return two Q-values Q1(s,a) and Q2(s,a) as outputs
        current_Q1, current_Q2 = self.critic(state, action)

```

```

    # Step 11: We compute the loss coming from the two Critic models: Critic_
    ↪ Loss = MSE_Loss(Q1(s,a), Qt) + MSE_Loss(Q2(s,a), Qt)
    critic_loss = F.mse_loss(current_Q1, target_Q) + F.mse_loss(current_Q2,
    ↪ target_Q)

    # Step 12: We backpropagate this Critic loss and update the parameters of
    ↪ the two Critic models with a SGD optimizer
    self.critic_optimizer.zero_grad()
    critic_loss.backward()
    self.critic_optimizer.step()

    # Step 13: Once every two iterations, we update our Actor model by
    ↪ performing gradient ascent on the output of the first Critic model
    if it % policy_freq == 0:
        actor_loss = -self.critic.Q1(state, self.actor(state)).mean()
        self.actor_optimizer.zero_grad()
        actor_loss.backward()
        self.actor_optimizer.step()

    # Step 14: Still once every two iterations, we update the weights of
    ↪ the Actor target by polyak averaging
    for param, target_param in zip(self.actor.parameters(), self.
    ↪ actor_target.parameters()):
        target_param.data.copy_(tau * param.data + (1 - tau) * target_param.
    ↪ data)

    # Step 15: Still once every two iterations, we update the weights of
    ↪ the Critic target by polyak averaging
    for param, target_param in zip(self.critic.parameters(), self.
    ↪ critic_target.parameters()):
        target_param.data.copy_(tau * param.data + (1 - tau) * target_param.
    ↪ data)

    # Making a save method to save a trained model
    def save(self, filename, directory):
        torch.save(self.actor.state_dict(), '%s/%s_actor.pth' % (directory,
    ↪ filename))
        torch.save(self.critic.state_dict(), '%s/%s_critic.pth' % (directory,
    ↪ filename))

    # Making a load method to load a pre-trained model
    def load(self, filename, directory):
        self.actor.load_state_dict(torch.load('%s/%s_actor.pth' % (directory,
    ↪ filename)))

```

```
self.critic.load_state_dict(torch.load('%s/%s_critic.pth' % (directory,
↪filename)))
```

1.7 We make a function that evaluates the policy by calculating its average reward over 10 episodes

```
[0]: def evaluate_policy(policy, eval_episodes=10):
    avg_reward = 0.
    for _ in range(eval_episodes):
        obs = env.reset()
        done = False
        while not done:
            action = policy.select_action(np.array(obs))
            obs, reward, done, _ = env.step(action)
            avg_reward += reward
        avg_reward /= eval_episodes
    print ("-----")
    print ("Average Reward over the Evaluation Step: %f" % (avg_reward))
    print ("-----")
    return avg_reward
```

1.8 We set the parameters

```
[0]: env_name = "AntBulletEnv-v0" # Name of a environment (set it to any Continuous
↪environment you want)
seed = 0 # Random seed number
start_timesteps = 1e4 # Number of iterations/timesteps before which the model
↪randomly chooses an action, and after which it starts to use the policy
↪network
eval_freq = 5e3 # How often the evaluation step is performed (after how many
↪timesteps)
max_timesteps = 5e5 # Total number of iterations/timesteps
save_models = True # Boolean checker whether or not to save the pre-trained
↪model
expl_noise = 0.1 # Exploration noise - STD value of exploration Gaussian noise
batch_size = 100 # Size of the batch
discount = 0.99 # Discount factor gamma, used in the calculation of the total
↪discounted reward
tau = 0.005 # Target network update rate
policy_noise = 0.2 # STD of Gaussian noise added to the actions for the
↪exploration purposes
noise_clip = 0.5 # Maximum value of the Gaussian noise added to the actions
↪(policy)
```

```
policy_freq = 2 # Number of iterations to wait before the policy network (Actor_
↪model) is updated
```

1.9 We create a file name for the two saved models: the Actor and Critic models

```
[0]: file_name = "%s_%s_%s" % ("TD3", env_name, str(seed))
print ("-----")
print ("Settings: %s" % (file_name))
print ("-----")
```

1.10 We create a folder inside which will be saved the trained models

```
[0]: if not os.path.exists("./results"):
    os.makedirs("./results")
if save_models and not os.path.exists("./pytorch_models"):
    os.makedirs("./pytorch_models")
```

1.11 We create the PyBullet environment

```
[0]: env = gym.make(env_name)
```

1.12 We set seeds and we get the necessary information on the states and actions in the chosen environment

```
[0]: env.seed(seed)
torch.manual_seed(seed)
np.random.seed(seed)
state_dim = env.observation_space.shape[0]
action_dim = env.action_space.shape[0]
max_action = float(env.action_space.high[0])
```

1.13 We create the policy network (the Actor model)

```
[0]: policy = TD3(state_dim, action_dim, max_action)
```

1.14 We create the Experience Replay memory

```
[0]: replay_buffer = ReplayBuffer()
```

1.15 We define a list where all the evaluation results over 10 episodes are stored

```
[0]: evaluations = [evaluate_policy(policy)]
```

1.16 We create a new folder directory in which the final results (videos of the agent) will be populated

```
[0]: def mkdir(base, name):
    path = os.path.join(base, name)
    if not os.path.exists(path):
        os.makedirs(path)
    return path
work_dir = mkdir('exp', 'brs')
monitor_dir = mkdir(work_dir, 'monitor')
max_episode_steps = env._max_episode_steps
save_env_vid = False
if save_env_vid:
    env = wrappers.Monitor(env, monitor_dir, force = True)
    env.reset()
```

1.17 We initialize the variables

```
[0]: total_timesteps = 0
timesteps_since_eval = 0
episode_num = 0
done = True
t0 = time.time()
```

1.18 Training

```
[0]: # We start the main loop over 500,000 timesteps
while total_timesteps < max_timesteps:

    # If the episode is done
    if done:

        # If we are not at the very beginning, we start the training process of the
        ↪model
        if total_timesteps != 0:
            print("Total Timesteps: {} Episode Num: {} Reward: {}".
            ↪format(total_timesteps, episode_num, episode_reward))
            policy.train(replay_buffer, episode_timesteps, batch_size, discount, tau,
            ↪policy_noise, noise_clip, policy_freq)
```



```

# We evaluate the episode and we save the policy
if timesteps_since_eval >= eval_freq:
    timesteps_since_eval %= eval_freq
    evaluations.append(evaluate_policy(policy))
    policy.save(file_name, directory="./pytorch_models")
    np.save("./results/%s" % (file_name), evaluations)

# When the training step is done, we reset the state of the environment
obs = env.reset()

# Set the Done to False
done = False

# Set rewards and episode timesteps to zero
episode_reward = 0
episode_timesteps = 0
episode_num += 1

# Before 10000 timesteps, we play random actions
if total_timesteps < start_timesteps:
    action = env.action_space.sample()
else: # After 10000 timesteps, we switch to the model
    action = policy.select_action(np.array(obs))
    # If the explore_noise parameter is not 0, we add noise to the action and
    ↪we clip it
    if expl_noise != 0:
        action = (action + np.random.normal(0, expl_noise, size=env.action_space.
        ↪shape[0])).clip(env.action_space.low, env.action_space.high)

# The agent performs the action in the environment, then reaches the next
↪state and receives the reward
new_obs, reward, done, _ = env.step(action)

# We check if the episode is done
done_bool = 0 if episode_timesteps + 1 == env._max_episode_steps else
↪float(done)

# We increase the total reward
episode_reward += reward

# We store the new transition into the Experience Replay memory (ReplayBuffer)
replay_buffer.add((obs, new_obs, action, reward, done_bool))

# We update the state, the episode timestep, the total timesteps, and the
↪timesteps since the evaluation of the policy
obs = new_obs

```

```

episode_timesteps += 1
total_timesteps += 1
timesteps_since_eval += 1

# We add the last policy evaluation to our list of evaluations and we save our
↪model
evaluations.append(evaluate_policy(policy))
if save_models: policy.save("%s" % (file_name), directory="./pytorch_models")
np.save("./results/%s" % (file_name), evaluations)

```

1.19 Inference

```

[0]: class Actor(nn.Module):

    def __init__(self, state_dim, action_dim, max_action):
        super(Actor, self).__init__()
        self.layer_1 = nn.Linear(state_dim, 400)
        self.layer_2 = nn.Linear(400, 300)
        self.layer_3 = nn.Linear(300, action_dim)
        self.max_action = max_action

    def forward(self, x):
        x = F.relu(self.layer_1(x))
        x = F.relu(self.layer_2(x))
        x = self.max_action * torch.tanh(self.layer_3(x))
        return x

class Critic(nn.Module):

    def __init__(self, state_dim, action_dim):
        super(Critic, self).__init__()
        # Defining the first Critic neural network
        self.layer_1 = nn.Linear(state_dim + action_dim, 400)
        self.layer_2 = nn.Linear(400, 300)
        self.layer_3 = nn.Linear(300, 1)
        # Defining the second Critic neural network
        self.layer_4 = nn.Linear(state_dim + action_dim, 400)
        self.layer_5 = nn.Linear(400, 300)
        self.layer_6 = nn.Linear(300, 1)

    def forward(self, x, u):
        xu = torch.cat([x, u], 1)
        # Forward-Propagation on the first Critic Neural Network
        x1 = F.relu(self.layer_1(xu))
        x1 = F.relu(self.layer_2(x1))
        x1 = self.layer_3(x1)

```

```

    # Forward-Propagation on the second Critic Neural Network
    x2 = F.relu(self.layer_4(xu))
    x2 = F.relu(self.layer_5(x2))
    x2 = self.layer_6(x2)
    return x1, x2

def Q1(self, x, u):
    xu = torch.cat([x, u], 1)
    x1 = F.relu(self.layer_1(xu))
    x1 = F.relu(self.layer_2(x1))
    x1 = self.layer_3(x1)
    return x1

# Selecting the device (CPU or GPU)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Building the whole Training Process into a class

class TD3(object):

    def __init__(self, state_dim, action_dim, max_action):
        self.actor = Actor(state_dim, action_dim, max_action).to(device)
        self.actor_target = Actor(state_dim, action_dim, max_action).to(device)
        self.actor_target.load_state_dict(self.actor.state_dict())
        self.actor_optimizer = torch.optim.Adam(self.actor.parameters())
        self.critic = Critic(state_dim, action_dim).to(device)
        self.critic_target = Critic(state_dim, action_dim).to(device)
        self.critic_target.load_state_dict(self.critic.state_dict())
        self.critic_optimizer = torch.optim.Adam(self.critic.parameters())
        self.max_action = max_action

    def select_action(self, state):
        state = torch.Tensor(state.reshape(1, -1)).to(device)
        return self.actor(state).cpu().data.numpy().flatten()

    def train(self, replay_buffer, iterations, batch_size=100, discount=0.99,
    ↪tau=0.005, policy_noise=0.2, noise_clip=0.5, policy_freq=2):

        for it in range(iterations):

            # Step 4: We sample a batch of transitions (s, s', a, r) from the memory
            batch_states, batch_next_states, batch_actions, batch_rewards,
            ↪batch_dones = replay_buffer.sample(batch_size)
            state = torch.Tensor(batch_states).to(device)
            next_state = torch.Tensor(batch_next_states).to(device)
            action = torch.Tensor(batch_actions).to(device)
            reward = torch.Tensor(batch_rewards).to(device)

```

```

done = torch.Tensor(batch_dones).to(device)

# Step 5: From the next state s', the Actor target plays the next action
→ a'
next_action = self.actor_target(next_state)

# Step 6: We add Gaussian noise to this next action a' and we clamp it in
→ a range of values supported by the environment
noise = torch.Tensor(batch_actions).data.normal_(0, policy_noise).
→ to(device)
noise = noise.clamp(-noise_clip, noise_clip)
next_action = (next_action + noise).clamp(-self.max_action, self.
→ max_action)

# Step 7: The two Critic targets take each the couple (s', a') as input
→ and return two Q-values Qt1(s',a') and Qt2(s',a') as outputs
target_Q1, target_Q2 = self.critic_target(next_state, next_action)

# Step 8: We keep the minimum of these two Q-values: min(Qt1, Qt2)
target_Q = torch.min(target_Q1, target_Q2)

# Step 9: We get the final target of the two Critic models, which is: Qt
→ = r + γ * min(Qt1, Qt2), where γ is the discount factor
target_Q = reward + ((1 - done) * discount * target_Q).detach()

# Step 10: The two Critic models take each the couple (s, a) as input and
→ return two Q-values Q1(s,a) and Q2(s,a) as outputs
current_Q1, current_Q2 = self.critic(state, action)

# Step 11: We compute the loss coming from the two Critic models: Critic
→ Loss = MSE_Loss(Q1(s,a), Qt) + MSE_Loss(Q2(s,a), Qt)
critic_loss = F.mse_loss(current_Q1, target_Q) + F.mse_loss(current_Q2,
→ target_Q)

# Step 12: We backpropagate this Critic loss and update the parameters of
→ the two Critic models with a SGD optimizer
self.critic_optimizer.zero_grad()
critic_loss.backward()
self.critic_optimizer.step()

# Step 13: Once every two iterations, we update our Actor model by
→ performing gradient ascent on the output of the first Critic model
if it % policy_freq == 0:
    actor_loss = -self.critic.Q1(state, self.actor(state)).mean()
    self.actor_optimizer.zero_grad()
    actor_loss.backward()

```

```

        self.actor_optimizer.step()

        # Step 14: Still once every two iterations, we update the weights of
        → the Actor target by polyak averaging
        for param, target_param in zip(self.critic.parameters(), self.
        → critic_target.parameters()):
            target_param.data.copy_(tau * param.data + (1 - tau) * target_param.
            → data)

        # Step 15: Still once every two iterations, we update the weights of
        → the Critic target by polyak averaging
        for param, target_param in zip(self.actor.parameters(), self.
        → actor_target.parameters()):
            target_param.data.copy_(tau * param.data + (1 - tau) * target_param.
            → data)

        # Making a save method to save a trained model
        def save(self, filename, directory):
            torch.save(self.actor.state_dict(), '%s/%s_actor.pth' % (directory,
            → filename))
            torch.save(self.critic.state_dict(), '%s/%s_critic.pth' % (directory,
            → filename))

        # Making a load method to load a pre-trained model
        def load(self, filename, directory):
            self.actor.load_state_dict(torch.load('%s/%s_actor.pth' % (directory,
            → filename)))
            self.critic.load_state_dict(torch.load('%s/%s_critic.pth' % (directory,
            → filename)))

def evaluate_policy(policy, eval_episodes=10):
    avg_reward = 0.
    for _ in range(eval_episodes):
        obs = env.reset()
        done = False
        while not done:
            action = policy.select_action(np.array(obs))
            obs, reward, done, _ = env.step(action)
            avg_reward += reward
    avg_reward /= eval_episodes
    print ("-----")
    print ("Average Reward over the Evaluation Step: %f" % (avg_reward))
    print ("-----")
    return avg_reward

env_name = "AntBulletEnv-v0"

```

```

seed = 0

file_name = "%s_%s_%s" % ("TD3", env_name, str(seed))
print ("-----")
print ("Settings: %s" % (file_name))
print ("-----")

eval_episodes = 10
save_env_vid = True
env = gym.make(env_name)
max_episode_steps = env._max_episode_steps
if save_env_vid:
    env = wrappers.Monitor(env, monitor_dir, force = True)
    env.reset()
env.seed(seed)
torch.manual_seed(seed)
np.random.seed(seed)
state_dim = env.observation_space.shape[0]
action_dim = env.action_space.shape[0]
max_action = float(env.action_space.high[0])
policy = TD3(state_dim, action_dim, max_action)
policy.load(file_name, './pytorch_models/')
_ = evaluate_policy(policy, eval_episodes=eval_episodes)

```